WIGVO: Real-Time Bidirectional Speech Translation over Legacy PSTN Calls via Dual-Session Echo Gating

Hyeong-seob Kim* **Sang-woo Son** **Hyun-woo Cho** **Hyeon-sang Kim** **Jin-mo Kim**
SoundMind Inc.
hskim2@soundmind.life

Abstract

Real-time speech translation with large language models (LLMs) has become feasible in controlled settings such as mobile apps and web browsers, where developers assume wideband audio and client-side echo cancellation. Recent end-to-end full-duplex models further push latency below 200 ms but remain confined to broadband environments. However, deploying such systems over the Public Switched Telephone Network (PSTN) remains challenging due to narrowband G.711 audio, unpredictable round-trip delays, and absence of client-side signal processing. We present **WIGVO** (WIGTN Voice-Only), a server-side relay system that enables bidirectional LLM-based speech translation over ordinary telephone calls without requiring app installation or carrier integration. A central contribution is addressing what we term *echo-induced self-reinforcing translation loops*: synthesized speech echoing back through the PSTN gets re-ingested and repeatedly translated. WIGVO solves this through a dual-session architecture with deterministic silence injection and energy-based voice activity detection (VAD) gating. We evaluate WIGVO on 155 PSTN calls (148 instrumented, 147 completed) across three communication modes—voice-to-voice (V2V), text-to-voice (T2V), and full-agent—observing 555 ms median caller-to-callee latency and 2,684 ms median callee-to-caller latency, zero echo-induced translation loops, and USD 0.28 per minute cost. The system is deployed at https://wigvo-web-gzjzn35jyq-du.a.run.app with a live demo and video walkthrough.

1 Introduction

Streaming speech-to-speech translation has advanced rapidly (Seamless Communication et al., 2023a; OpenAI, 2024), but existing systems assume wideband audio and client-side acoustic echo cancellation (AEC) (ITU-T, 2015). The Public Switched Telephone Network (PSTN) remains the primary inbound interface for local businesses, hospitals, and government offices (International Telecommunication Union, 2020), yet it encodes audio via G.711 μ-law at 8 kHz (ITU-T, 1988), introduces 80–600 ms delays (ITU-T, 2003), and provides no AEC. For foreign residents, people with speech anxiety, or users with hearing impairments, making such calls remains a significant barrier.

We present **WIGVO**, a server-side relay that bridges a web client and a standard PSTN phone number through two concurrent LLM-backed streaming sessions. The caller speaks (or types) via a browser; the callee answers on an ordinary phone—no app installation, carrier integration, or special hardware is required on either end. WIGVO translates bidirectionally in real time, supporting voice-to-voice (V2V), text-to-voice (T2V), and full-agent modes.

A key challenge is what we call an **echo-induced self-reinforcing translation loop**: synthesized speech echoing back through the telephone network gets re-recognized and repeatedly translated, creating a runaway feedback cycle (Section 3.1). Our solution combines three architectural mechanisms—directional separation, deterministic silence injection, and energy-based gating—described in Section 3.2.

We evaluate WIGVO on 147 completed PSTN calls and report latency, echo suppression, and cost metrics. Our contributions are:

- We formalize the *echo-induced translation loop* problem in streaming speech translation over telephony.
- We propose a *dual-session gated architecture* combining directional separation, silence injection, and energy-based gating for PSTN environments.

*Corresponding author. Authors 3–5 listed alphabetically.

- We deploy and evaluate a working relay server across three communication modes, achieving zero echo-induced loops across 147 completed PSTN calls with 555 ms median caller-to-callee latency at USD 0.28 per minute.

The modular cascade architecture enables language extensibility through component substitution, and the dual-session gated design generalizes to multilingual call centers, voice inclusion services, and telephony-native LLM agents.

2 Related Work

Simultaneous speech translation systems such as SeamlessM4T (Seamless Communication et al., 2023b), Seamless Streaming (Seamless Communication et al., 2023a), and ESPnet-ST-v2 (Yan et al., 2023) achieve real-time quality but assume wideband audio from controlled environments. End-to-end full-duplex models—Moshi (Défossez et al., 2024) (200 ms latency), PersonaPlex (Roy et al., 2026) (70 ms turn-taking), and Hibiki (Labiausse et al., 2025) (near-human-interpreter quality, French→English)—all require clean 24 kHz audio and do not address G.711 codec artifacts, telephony echo, or narrowband VAD failures. In telephony AI, Google Duplex (Leviathan and Matias, 2018) performs monolingual task completion over PSTN, while Vapi[1] and Bland.ai[2] provide LLM-powered phone agents without cross-lingual translation. For accessibility, FCC Telecommunications Relay Services (Federal Communications Commission, 2024) require human intermediaries limited to specific countries, while sign.mt (Moryossef, 2024) targets sign language rather than telephony. Table 1 positions WIGVO against these systems.

3 System Architecture and Echo Gating

Figure 1 shows the high-level architecture. A browser client connects to the relay server via WebSocket, sending 16 kHz PCM audio. The relay manages two independent Realtime LLM sessions and a Twilio telephony gateway (Twilio, 2024). An `AudioRouter` dispatches events to one of three pipelines—V2V (streaming ASR+TTS in both directions), T2V (typed text to TTS for the callee, speech-to-captions for the caller), and Full-Agent (autonomous call management)—via the Strategy pattern.

[1] https://vapi.ai
[2] https://bland.ai

System	PSTN	Bidir.	S2S	Echo	A11y
SeamlessM4T	×	✓	✓	N/A	×
Moshi / Hibiki	×	×	✓	N/A	×
Google Duplex	✓	×	×	N/D	×
FCC TRS	✓	✓	×	Human	✓
sign.mt	×	✓	✓	N/A	✓
Vapi / Bland	✓	×	×	N/D	×
WIGVO	✓	✓	✓	✓	✓

Table 1: Comparison with existing systems. PSTN: works over telephone network; Bidir.: bidirectional translation; S2S: automated speech-to-speech (not human-mediated); Echo: handles telephony echo; A11y: accessibility modes. N/D: not disclosed.

Dual Realtime sessions. Session A receives browser audio (PCM16) and produces translated G.711 μ-law for the telephony gateway. Session B receives PSTN audio (G.711 μ-law, 8 kHz) and produces translated audio or captions for the browser. In V2V mode, Session B performs end-to-end translation with TTS output via the Realtime API; in T2V and Full-Agent modes, it performs STT only, delegating translation to a separate Chat Completion call (GPT-4o-mini, `temperature=0`) to prevent generative hallucination—the Realtime API's tendency to produce contextually plausible but untranslated additions. Each session maintains its own system prompt and sliding context window, ensuring strict directional separation.

3.1 The Echo Loop Problem

In a naive single-session design, TTS output played to the callee echoes back through the PSTN after 80–600 ms, gets re-recognized, and triggers another translation cycle—generating progressively distorted paraphrases until manually interrupted. A Pearson-correlation echo detector comparing outgoing TTS with incoming audio failed in production due to μ-law nonlinear quantization and variable delays. Without gating, eight of ten test calls produced echo loops; the correlation detector reduced this to three of ten but introduced frequent false positives. Echo Gate eliminated loops entirely (0 of 147 completed calls).

3.2 Echo Gate and 3-Stage Pipeline

PSTN audio entering Session B passes through a 3-stage filter pipeline (Figure 2):

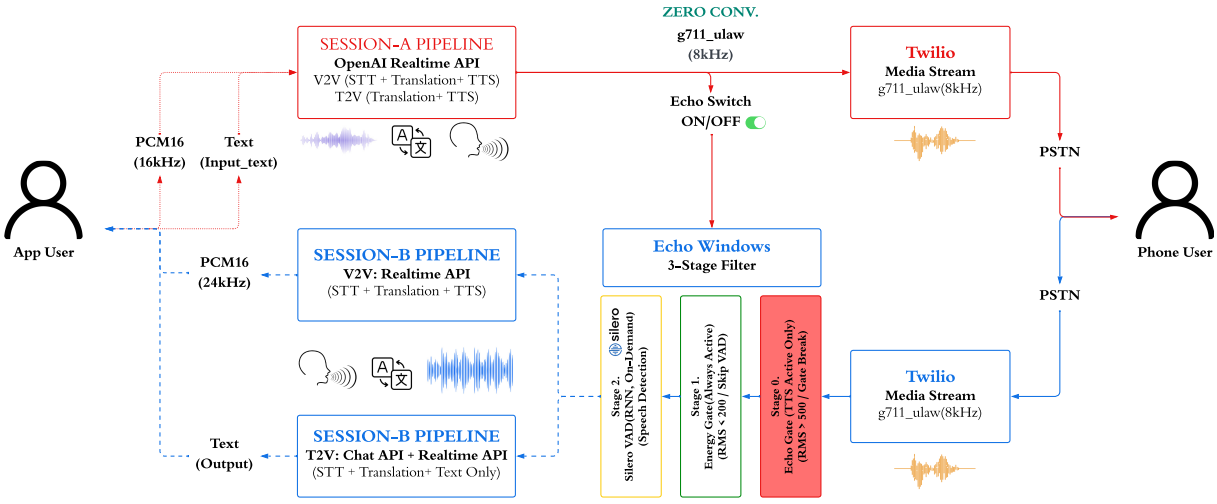

Figure 1: WIGVO system architecture. Session A (red) translates user speech and outputs G.711 directly to Twilio; Session B (blue) receives PSTN audio through a 3-stage filter pipeline (Echo Gate → Energy Gate → Silero VAD). A dashed control signal from Session A activates the Echo Gate in Session B's filter pipeline, preventing feedback loops by injecting μ-law silence during TTS playback.

Stage 0: Echo Gate. When Session A begins streaming TTS, an *echo window* is activated (duration estimated from byte length, capped at 1.2 s). Incoming PSTN audio is replaced with μ-law silence (0xFF) before forwarding to Session B. A 0.3 s dynamic cooldown (0.5 s for T2V) accounts for PSTN jitter, followed by a dynamic post-echo settling window (TTS duration $\times$ 0.3, clamped to 0.5–1.5 s) that suppresses AGC recovery noise while permitting high-energy breakthrough ($\geq$400 RMS) for genuine callee speech.

Stage 1: RMS Energy Gate. During echo windows (Figure 2B), only signals $\geq$400 RMS break through as genuine speech; typical echo energy (100–400 RMS) remains suppressed. Outside echo windows, a 150 RMS threshold filters line noise. A three-level interrupt priority (callee $>$ caller $>$ AI TTS) enables natural turn-taking.

Stage 2: Silero VAD. Frames passing the energy gate are processed by Silero VAD (Silero Team, 2021). PSTN audio (8 kHz) is upsampled to 16 kHz via zero-order hold before inference. Asymmetric hysteresis—onset at 96 ms (probability $\geq$ 0.5, three frames), offset at 480 ms (probability $<$ 0.35, 15 frames)—reduces `speech_stopped` latency from 15–72 s (server VAD on PSTN) to 480 ms.

Silence injection in Stage 0 addresses a distinct failure mode beyond comfort noise generation (Zopf, 2002): if the relay simply *drops* audio during echo windows, the Realtime API's server-side VAD never observes a silence-to-speech transition and remains stuck in a "speaking" state indefinitely. To our knowledge, this failure mode and its mitigation have not been reported in neural streaming S2ST pipelines.

3.3 PSTN VAD Failures and Robustness

The Realtime API's server-side VAD works well for wideband WebRTC audio but fails on the telephone network in three ways: (1) constant low-level noise (line hum, codec artifacts) is classified as speech, causing stuck durations of 15–72 s; (2) the noisy telephone floor prevents energy from dropping below the silence threshold, delaying end-of-utterance detection; (3) audio discontinuities during echo suppression prevent clean silence detection. The three-stage pipeline addresses all three: the echo gate eliminates discontinuities, the energy gate removes sub-speech noise, and Silero VAD provides neural speech detection with hysteresis. A 15-pattern STT hallucination blocklist intercepts broadcast-style Whisper (Radford et al., 2023) artifacts; the blocklist intercepted 100 hallucinations across 148 instrumented calls.

Across 147 completed calls, the echo gate activated 1,046 times with zero echo-induced loops.

4 Evaluation

We evaluated WIGVO on 155 PSTN calls (148 fully instrumented) for Korean↔English translation over a five-day period (Feb 23–27, 2026). Of these, 147 calls completed successfully; three

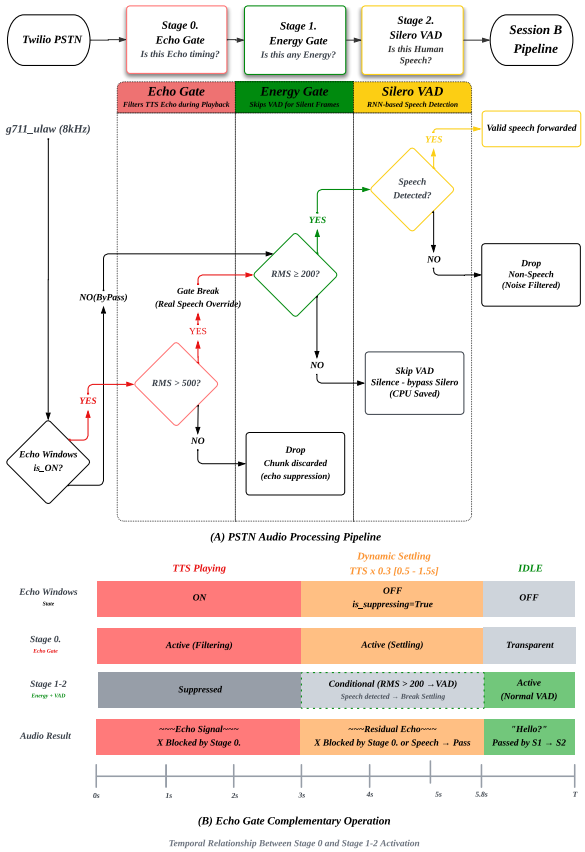

Figure 2: PSTN audio processing pipeline. (A) Three-stage filter between Twilio input and Session B. (B) Temporal behavior: during TTS playback, the Echo Gate replaces audio with μ-law silence while Stages 1–2 are suppressed; after settling, Stage 0 becomes transparent and Stages 1–2 filter PSTN noise.

failed and five were in-progress at collection time.

4.1 Latency

Table 2 reports latency statistics across the 148 instrumented calls.

Session A achieves 555 ms median latency, within the range for interactive communication. Session B, which depends on PSTN audio quality and streaming ASR, shows 2,684 ms median. STT (Whisper) dominates Session B latency: across 525 paired turns, mean STT latency (3,544 ms) accounts for the majority of end-to-end time. Figure 3 shows the end-to-end latency distributions for Session A and Session B. The P95 (9,963 ms) is driven primarily by PSTN VAD stuck events—server-side VAD on noisy telephone audio occasionally delays `speech_stopped` detection to 15–16 s, independent of utterance length.

T2V mode shows lower Session A latency (no caller-side ASR), while Agent mode shows higher Session B latency due to function-calling overhead.

Stage	Avg	P50	P95	Max	N
Session A (ms)	619	555	1,169	3,906	739
Session B E2E (ms)	3,650	2,684	9,963	18,234	683
STT (ms)	3,544	2,601	9,392	—	525*
Translation (ms)	535	84	1,961	—	525*
First message (ms)	2,081	1,248	5,751	—	148

Table 2: Translation latency across 148 instrumented calls (Feb 23–27). Session A: caller→callee (ASR+translate+TTS). Session B: callee→caller, decomposed into STT (Whisper) and translation components. N varies by metric scope; *STT and Translation N=525 reflects paired turns with both measurements (139 misaligned turns skipped). Translation latency reflects Chat API (GPT-4o-mini) for T2V/Agent modes and Realtime API for V2V mode.

While Session B's ∼2.7 s median exceeds ITU-T G.114's 150 ms one-way threshold (ITU-T, 2003), that standard applies to same-language dialogue. A more appropriate baseline is professional simultaneous interpretation, where ear-voice span ranges from 2 to 5 s (Goldman-Eisler, 1972; Timarová et al., 2011); Session B falls at the lower bound. The latency is structurally asymmetric by design: Session A (555 ms) keeps the caller-to-callee direction interactive, while Session B prioritizes STT accuracy on degraded G.711 audio. The bottleneck is narrowband ASR, not the relay architecture; replacing Whisper with a streaming STT model (Kyutai, 2025) could reduce Session B below 800 ms.

Figure 4 shows the correlation between utterance length and Session B latency ($r = 0.333$, $p < 0.001$), confirming that VAD stuck events are independent of utterance length.

4.2 Echo and Safety

Across the 148 instrumented calls, the echo gate activated 1,046 times (7.1/call), corresponding to TTS playback events. Critically, we observed zero echo-induced translation loops, whereas early prototypes without gating reliably produced such loops. The echo gate permitted 354 callee interruptions during active TTS playback (gate break), confirming that the energy-based override preserves natural turn-taking. Additionally, 17 settling breakthroughs demonstrate that the conditional VAD activation (Section 3.2) successfully detects genuine speech during the post-echo set-

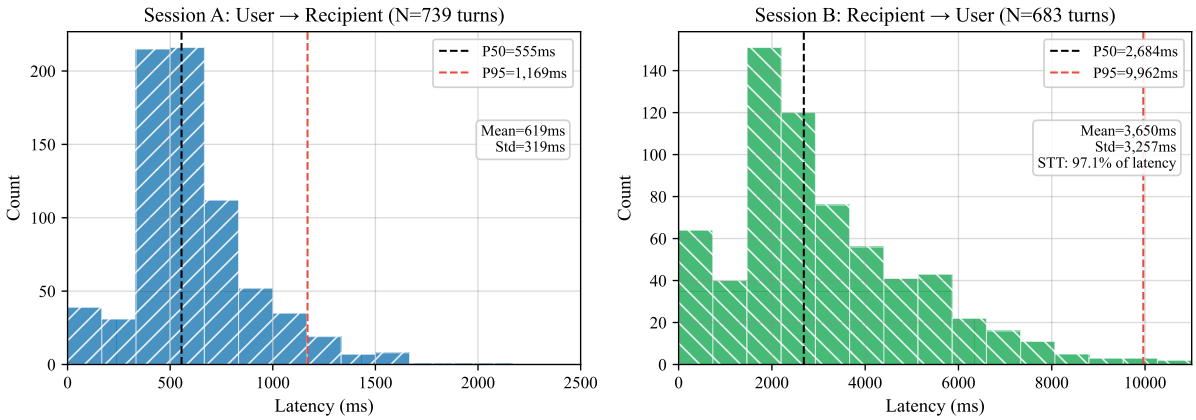

Figure 3: End-to-end latency distributions for Session A (caller→callee, N=739 turns) and Session B (callee→caller, N=683 turns) over live PSTN calls.

tling period. A minimum speech duration filter (480 ms hysteresis) mitigated VAD false triggers to 277 total (1.9/call). The guardrail module blocked 100 hallucinated outputs (0.68/call), including broadcast-style phrases from Whisper on low-energy input and structural noise patterns (consonant-only sequences, repetitive tokens).

4.3 Cost

Across the 148 instrumented calls (220.2 minutes total), WIGVO consumed 1,870,657 tokens at a total cost of USD 61.26, yielding USD 0.28/min (USD 0.41/call). These costs reflect dual-session Realtime API pricing and remain an order of magnitude cheaper than professional telephone interpretation services (USD 1–3/min).[3] Self-hosted STT alternatives could further reduce per-minute cost; a formal user study with target populations (Brooke, 1996) is planned as future work.

4.4 Translation Quality

We evaluated translation quality by comparing WIGVO's real-time outputs against independent offline references generated by Gemini 2.5 Flash, a separate model family from the GPT-4o Realtime API used for live translation.[4] This cross-model design avoids same-model evaluation bias. From 155 calls with bilingual transcripts, we extracted 215 caller-to-callee (en→ko) and 558 callee-to-caller (ko→en) segments. A noise filter removed 40 segments (5.2%) containing non-linguistic inputs (consonant-only keyboard artifacts, repeated onomatopoeia, background TV transcriptions).

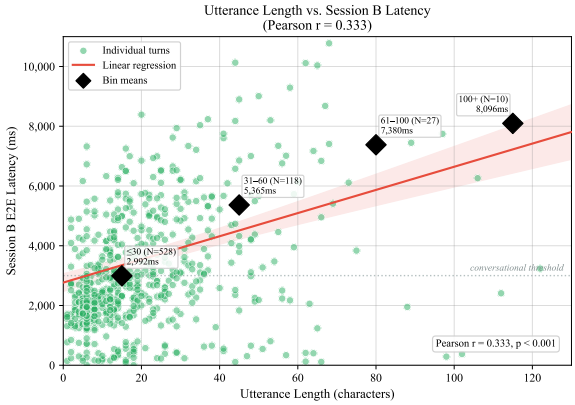

Figure 4: Utterance length (characters) vs. Session B E2E latency. Pearson $r = 0.333$ ($p < 0.001$), indicating that longer recipient utterances incur higher ASR-dominated latency.

Table 3 reports BLEU (Papineni et al., 2002), chrF2++ (Popović, 2015), and COMET (Rei et al., 2020) scores computed with SacreBLEU (Post, 2018) and `wmt22-comet-da`.

Surface-level n-gram metrics (BLEU, chrF2++) remain modest because real-time translation inherently produces valid paraphrases that differ lexically from offline references—a known limitation of these metrics for simultaneous interpretation. COMET, which evaluates semantic adequacy via cross-lingual embeddings, provides a more mean-

[3] Based on published per-minute rates for over-the-phone interpretation (OPI) from major commercial providers.

[4] In T2V and Agent modes, Session B translation uses GPT-4o-mini Chat Completion (`temperature=0`); V2V mode retains the Realtime API pipeline (Section 3).

Direction	N	BLEU	chrF2++	COMET
en→ko	212	15.9	31.7	0.7078
ko→en	521	17.4	27.8	0.6242

Table 3: Translation quality against Gemini 2.5 Flash offline references. COMET (`wmt22-comet-da`) is a neural metric trained on human judgments; higher is better (max 1.0).

ingful signal: en→ko achieves 0.71, indicating adequate translation quality despite surface divergence. The lower ko→en score (0.62) reflects STT errors on narrowband telephony audio propagating through the pipeline, consistent with the ASR bottleneck identified in the latency analysis above.

4.5 Baseline Comparison

Table 4 contextualizes WIGVO's latency against related systems. No existing system performs bidirectional cross-lingual speech translation over PSTN, so we report absolute metrics rather than head-to-head comparisons.

Twilio's industry benchmark for AI voice agents targets 1,115 ms median turn gap (Twilio, 2025). Session A (555 ms) falls well within this target; Session B (2,684 ms) exceeds it but includes cross-lingual translation that monolingual agents do not perform.

5 Demonstration

WIGVO is deployed at `https://wigvo-web-gzjzn35jyq-du.a.run.app`.[5] The source code is available at `https://github.com/wigtn/wigvo-v2` under the MIT License. A reviewer session proceeds as: (1) select a scenario, (2) chat with the AI agent to provide caller intent and phone number, (3) the system places a PSTN call with bidirectional translation, and (4) real-time captions appear alongside call status. Mode switching (V2V, T2V, Full-Agent) is available during active calls. Figure 5 shows the interface during a V2V call.

[5]Demo video: `https://youtu.be/PLACEHOLDER`. Reviewer login: `demo@gmail.com` / `wigvo2026!` (authentication required to manage API costs).

System	Network	Bidir.	Latency
Seamless Str.	Wideband	×	<2 s AL
Moshi	Wideband	×†	∼200 ms
Hibiki	Wideband	×	2.8–5.6 s LAAL
Google Duplex	PSTN	×‡	N/D
Vapi	WebRTC/PSTN	×‡	∼965 ms*
Bland.ai	PSTN	×‡	<400 ms*
WIGVO SA	**PSTN**	✓	**555 ms P50**
WIGVO SB	**PSTN**	✓	**2,684 ms P50**

Table 4: Latency comparison with related systems. AL: Average Lagging; LAAL: Length-Adaptive AL. †: full-duplex but monolingual (no translation). ‡: monolingual voice agent (no cross-lingual). ∗: vendor-reported; independent reviews report higher values. N/D: not disclosed. Bidir.: bidirectional cross-lingual translation.

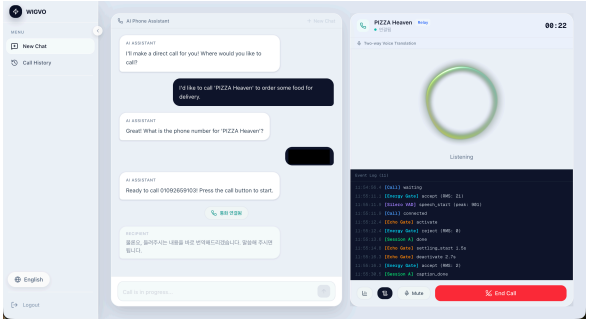

Figure 5: WIGVO web interface during a V2V call. Left: chat and caption view with bidirectional translations. Right: call status, duration, and mode indicator.

The current evaluation covers Korean–English only; the modular cascade architecture supports multilingual extension through component substitution.

6 Conclusion

We presented WIGVO, a dual-session streaming relay that enables bidirectional LLM-based speech translation over legacy PSTN calls. By combining deterministic silence injection with energy-based VAD gating, the system eliminates echo-induced translation loops—a failure mode we formalize for the first time in the context of streaming S2ST over telephony. Across 147 completed calls, the echo gate achieved a zero-loop rate while preserving natural turn-taking through 354 callee interruptions. Session A delivers 555 ms median latency, within interactive thresholds, at USD 0.28/min. The modular cascade architecture enables language extension through component substitution; replacing Whisper with a streaming telephony-optimized STT model is the most direct path to reducing Session B latency below 1 s. A formal user study with the target populations—foreign residents, individuals with speech anxiety, and people with hearing impairments—is planned as future work.

Ethics Statement

All calls were initiated by the authors and collaborators for system testing. Audio was logged for debugging and aggregate metrics only, with no personally identifiable information retained beyond temporary phone numbers. For deployment with real users, informed consent, data minimization, and compliance with privacy regulations (e.g., GDPR, PIPA) are required. The accessibility use case (T2V for speech anxiety) aims to reduce communication barriers but should be deployed with user agency and opt-in participation. We acknowledge the dual-use risk of automated voice translation over PSTN: the technology could be misused for voice phishing (vishing) or automated spam calls across language barriers. Mitigations include rate limiting, caller authentication, and abuse monitoring, which we plan to implement before public release. Translation errors in high-stakes contexts (medical, legal) could cause harm; the system currently provides no confidence indicators, and users should be informed that translations are AI-generated and not certified.

References

John Brooke. 1996. SUS: A 'quick and dirty' usability scale. *Usability Evaluation in Industry*, pages 189–194.

Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. 2024. Moshi: A speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*.

Federal Communications Commission. 2024. Telecommunications relay services (TRS). https://www.fcc.gov/consumers/guides/telecommunications-relay-service-trs.

Frieda Goldman-Eisler. 1972. Segmentation of input in simultaneous translation. *Journal of Psycholinguistic Research*, 1(2):127–140.

International Telecommunication Union. 2020. Public switched telephone network (PSTN): Technical overview. ITU-T Recommendation Series Q and G. https://www.itu.int/rec/T-REC-G/en.

ITU-T. 1988. Pulse code modulation (PCM) of voice frequencies. Technical Report Recommendation G.711, International Telecommunication Union.

ITU-T. 2003. One-way transmission time. Technical Report Recommendation G.114, International Telecommunication Union. Defines 150 ms and 400 ms one-way delay thresholds for conversational speech.

ITU-T. 2015. Digital network echo cancellers. Technical Report Recommendation G.168, International Telecommunication Union.

Kyutai. 2025. Kyutai STT: A speech-to-text model optimized for real-time usage. https://kyutai.org/stt.

Tom Labiausse, Laurent Mazaré, Edouard Grave, Patrick Pérez, Alexandre Défossez, and Neil Zeghidour. 2025. High-fidelity simultaneous speech-to-speech translation. In *Proceedings of the 42nd International Conference on Machine Learning*, pages 32116–32129.

Yaniv Leviathan and Yossi Matias. 2018. Google Duplex: An AI system for accomplishing real-world tasks over the phone. Google AI Blog, May 2018. https://ai.googleblog.com/2018/05/duplex-ai-system-for-natural-conversation.html.

Amit Moryossef. 2024. sign.mt: Real-time multilingual sign language translation application. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 182–190. Association for Computational Linguistics.

OpenAI. 2024. OpenAI realtime API documentation. https://platform.openai.com/docs/guides/realtime.

Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. BLEU: A method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318. Association for Computational Linguistics.

Maja Popović. 2015. chrF: Character n-gram F-score for automatic MT evaluation. In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395. Association for Computational Linguistics.

Matt Post. 2018. A call for clarity in reporting BLEU scores. In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 186–191. Association for Computational Linguistics.

Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. 2023. Robust speech recognition via large-scale weak supervision. In *Proceedings of the 40th International Conference on Machine Learning*, pages 28492–28518. PMLR.

Ricardo Rei, Craig Stewart, Ana C Farinha, and Alon Lavie. 2020. COMET: A neural framework for MT evaluation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 2685–2702. Association for Computational Linguistics.

Rajarshi Roy, Jonathan Raiman, Sang-gil Lee, Teodor-Dumitru Ene, Robert Kirby, Sungwon Kim, Jaehyeon Kim, and Bryan Catanzaro. 2026. Personaplex: Voice and role control for full duplex conversational speech models. *Preprint*, arXiv:2602.06053.

Seamless Communication et al. 2023a. Seamless: Multilingual expressive and streaming speech translation. In *arXiv preprint arXiv:2312.05187*. https://arxiv.org/abs/2312.05187.

Seamless Communication et al. 2023b. SeamlessM4T: Massively multilingual & multimodal machine translation. In *arXiv preprint arXiv:2308.11596*. https://arxiv.org/abs/2308.11596.

Silero Team. 2021. Silero VAD: Pre-trained enterprise-grade voice activity detector. https://github.com/snakers4/silero-vad.

Šárka Timaróvá, Ivonne Cerdá, and Reine Meylaerts. 2011. The effect of directionality on processing time in simultaneous interpreting. *Across Languages and Cultures*, 12(2):235–254. Reports ear-voice span of 2–4 s across language pairs.

Twilio. 2024. Twilio media streams documentation. https://www.twilio.com/docs/voice/media-streams.

Twilio. 2025. A guide to core latency in AI voice agents. https://www.twilio.com/en-us/blog/developers/best-practices/guide-core-latency-ai-voice-agents.

Brian Yan, Jiatong Shi, Yun Tang, Hirofumi Inaguma, Yifan Peng, Siddhant Dalmia, Peter Polák, Patrick Fernandes, Dan Berrebbi, Tomoki Hayashi, Xiaohui Zhang, Zhaoheng Ni, Moto Hira, Soumi Maiti, Juan Pino, and Shinji Watanabe. 2023. ESPnet-ST-v2: Multipurpose spoken language translation toolkit. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 400–411. Association for Computational Linguistics.

R. Zopf. 2002. RTP payload for comfort noise (CN). Technical Report RFC 3389, IETF.